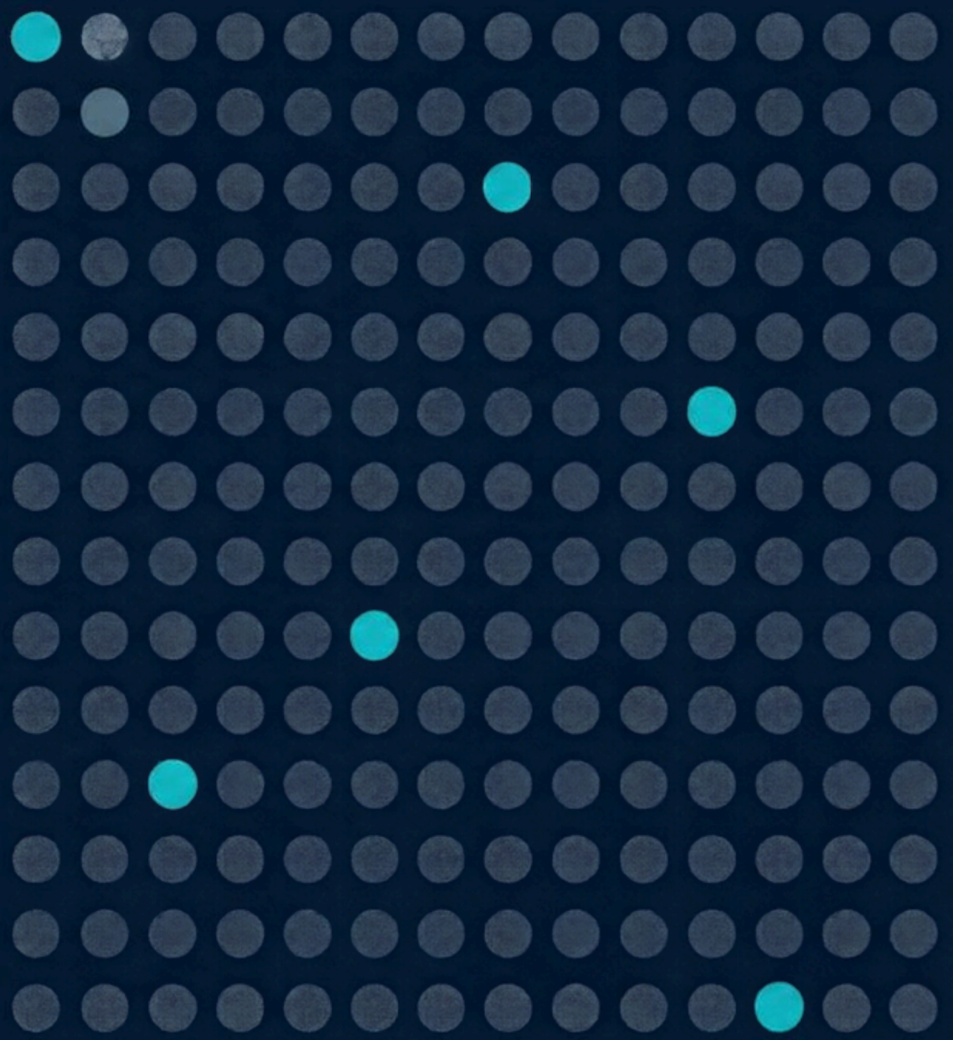
Fraud alerts freeze innocent people, and nobody counts them



Most AML flags are innocent customers

Frozen accounts mean
missed rent, payroll and bills

Banks tune for catching
crime, not for limiting harm

Set the harm cap first, then catch as much as we can

The bank sets a cap, for example 5 wrongly held innocents per 10,000.

1

Bank sets cap

→

2

System catches as much laundering as it can

→

3

Cap is tested on later data

Does the cap hold?

Tested on later time windows and a second dataset

At 5 per 10,000, the cap held in 6 of 8 windows on our main dataset (recall 32.8%)

And in 26 of 40 windows on a second dataset, nothing retuned (recall 13.4%)

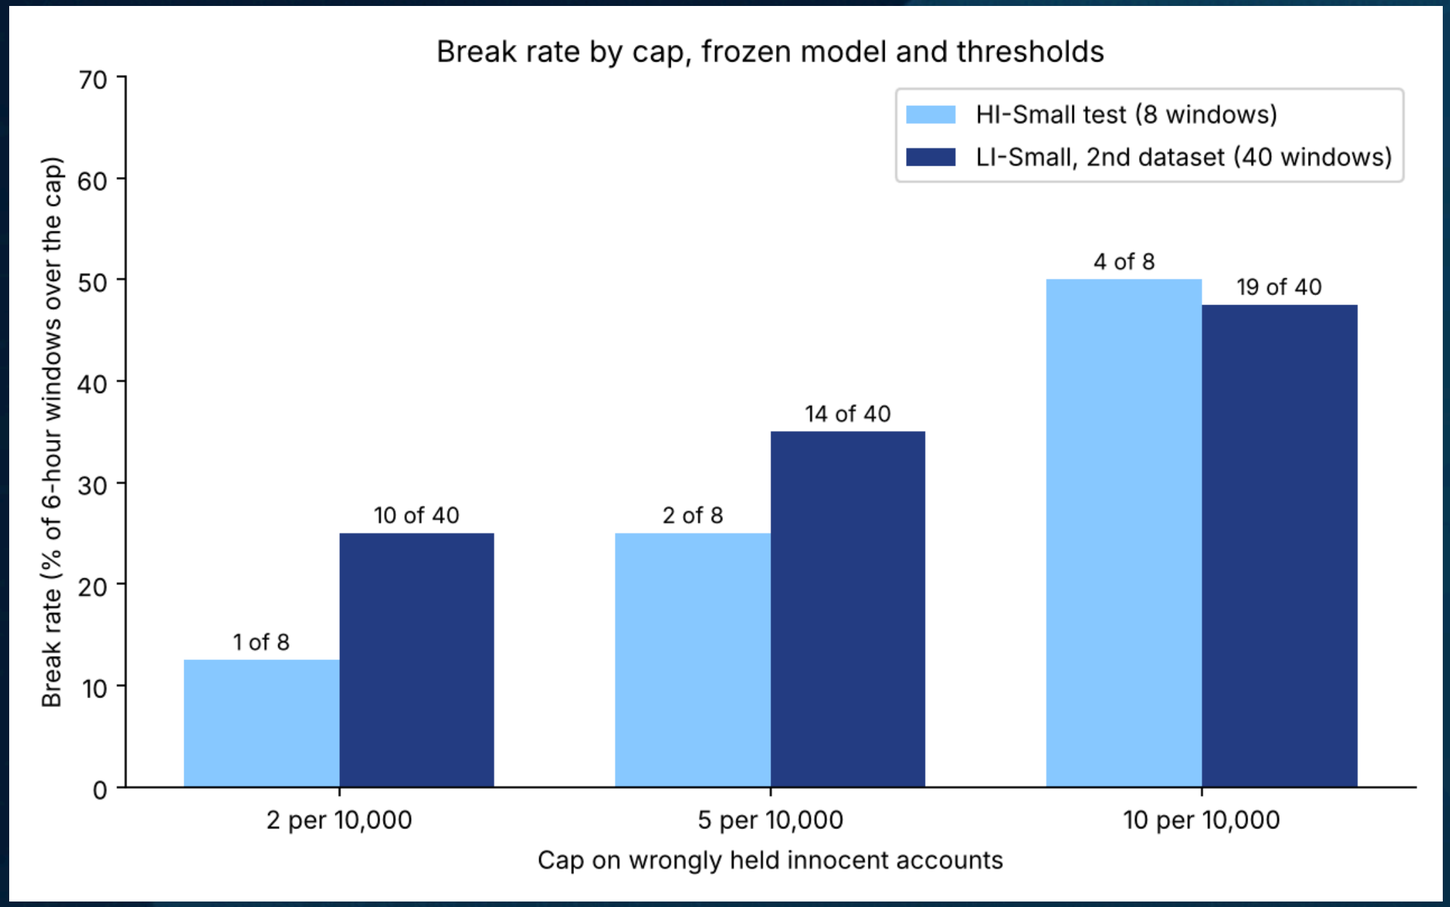


Chart goes here

Limitations

Synthetic data: "innocent" means "not labelled laundering"

Only 8 test windows on the main dataset, so "6 of 8" is coarse

Model and thresholds frozen on earlier data, not recalibrated

A proof of method, not a claim about a real bank

